

Flexible Scientific Data Management for Plant Phenomics Research

Peter Ansell¹, Robert Furbank², Kutila Gunasekera¹, Jianming Guo², David Benn³, Gareth Williams³, Xavier Sirault²

¹ eResearch Group, School of Information Technology and Electronic Engineering,
University of Queensland, Brisbane, Australia

² CSIRO Plant industry, High Resolution Plant Phenomics Centre, Canberra,
Australia

³ CSIRO IMT Advanced Scientific Computing and Research Data Services,
Melbourne, Australia

Abstract. In this paper, we expand on the design and implementation of the Phenomics Ontology Driven Data repository [1] (PODD) with respect to the capture, storage and retrieval of data and metadata generated at the High Resolution Plant Phenomics Centre (Canberra, Australia). PODD is a schema-driven Semantic Web database which uses the Resource Description Framework (RDF) model to store semi-structured information. RDF allows PODD to process information about a range of phenomics experiments without needing to define a universal schema for all of the different structures. To illustrate the process, exemplar datasets were generated using a medium throughput, high resolution, three-dimensional digitisation system purposely built for studying plant structure and function simultaneously under specific environmental conditions. The High Performance Compute (HPC), storage and data collection publication aspects of the workflow and their realisation in CSIRO infrastructure are also discussed along with their relationship to PODD.

1 Introduction

Since the genomics era, biology has become a data-driven science. Advances in robotics, automation and imaging, in combination with high performance computing have permitted the rapid production of large and complex biological datasets. Currently, high volumes of heterogeneous image data, physiological and morphological measurements are being acquired by a range of new phenotyping platforms located in purpose built phenomics centres across the world. These large datasets of phenotypic characters such as growth rate, plant architecture, photosynthetic performance, yield must be stored and related to genotype. This link is essential for expediting the identification of genetic variation in natural and derived genetic populations (e.g. germplasm collections, association genetic panels, recombinant inbred lines), understanding the dynamic relationship between phenotype, genotype and environment and delivering the increase in productivity necessary for feeding the world.

To achieve this goal, the vast array of phenotypic data collected from a variety of phenomics platforms must be combined with metadata explaining how the raw information was collected. This process forms structured digital objects containing the raw information. These objects are then delivered to a range of analysis pipelines, which transform the raw information into aggregated multi-phase datasets, each phase representing a new aggregation or inference from the original raw data. This reduction process converts the raw multi-dimensional data into information which is conceptually interpretable by a human being, i.e. new knowledge. The details about the steps taken are recorded to give context to the data. These links enable scientists to generate novel conclusions about the plants being phenotyped, including conclusions based on the genotype, plant growth environments, treatments, phenotyping platforms and the phenotyping and analysis processes.

To make sense of this large amount of information, sophisticated storage, archiving, searching and analysis capabilities are required. To date solutions to this problem have been handled essentially by private companies, and no suitable solution exists in the public domain. Lack of systems, both to manage linked metadata, and controlled vocabularies to describe plant growth and experimental conditions, have severely hampered sharing of plant phenomics data, comparison of results between laboratories and the capacity to carry out meta-analysis of existing data sets.

Thus, to support publicly-funded phenomics activities in Australia, the Phenomics Ontology Driven Data repository (PODD) has been developed as a repository for data produced by the variety of plant imaging and phenotyping platforms available at the High Resolution Plant Phenomics Centre, as well as for recording the contextual metadata associated with plant genotypes, treatments and environmental conditions [1]. PODD manages the links between the raw image files generated by plant imaging platforms, the results from analyses, and conclusions by scientists.

In this paper, we describe the workflow management that the High Resolution Plant Phenomics Centre (HRPPC) has implemented for keeping track of its phenomics data, metadata and experimental processes. This complex challenge was addressed by building a multi-disciplinary group of information technology experts and embedding users of phenomics technologies into it. The result of the approach is a state of the art computational and data mining environment, optimised for data access, data discovery and data sharing, which also provides the flexibility for linking genomic information through the use of RDF triples. In this context, we also describe the role of the CSIRO Data Access Portal (DAP) [2] to annotate and store raw and processed datasets. DAP also provides long term secure storage for data collections and the ability to search for, control access to, and cite them via Digital Object Identifiers. PODD manages the mapping of collections located in DAP to PODD projects, providing for the storage of large images and documents unsuited to RDF databases. Figure 1 shows the relationship between components and key data flows.



Fig. 1. HRPPC component relationships and data flow

2 Phenomics Ontology Driven Data repository

2.1 Semantic science for phenomics data management

Scientists have focused on including semantics into datasets, typically using the foundations of RDF and OWL, from two main directions. Some focus on defining ontologies based on hierarchies of scientific concepts and properties, while others have focused on mapping complex scientific datasets to RDF using syntax transformations without initially defining the semantic meaning of the results. In reality, most efforts fall somewhere in the middle, with ontological annotations attached to some data points while other nearby data points are syntactically represented using RDF, without links to ontologies of scientific concepts.

Increasingly however, providers of scientific datasets are focusing on enhancing their datasets using curated scientific concepts from ontologies. For example, scientists have used the Gene Ontology [3] to link well known concepts to represent common elements across genomics datasets. If these annotations are given using recognised URIs, and not as key-value string tuples, various RDF documents can be merged automatically based on RDF semantics to discover novel conclusions. Initially, different providers were not able to decide on URIs to use in Linked Data. This resulted in various third parties such as Bio2RDF [4] creating RDF versions of scientific datasets so that the RDF documents could be easily merged based on common URIs. However, recently more providers have given authoritative URIs for each object in their dataset, making it easy to integrate data from various locations.

2.2 Implementation of the Phenomics Ontology Driven Data repository

As described earlier, the PODD repository relies on semantic web technologies to manage phenomics data and metadata. Although both ontologies and mappings are essential, in PODD it was necessary to build the system with a relaxed ontological vocabulary to allow scientists to generate datasets before they were certain about the semantic meaning of their work. In addition, it was necessary for PODD to allow scientists to continue to maintain projects containing curated scientific concepts alongside raw experimental data.

The basic PODD project is structured as a tree, which, when used with PlantScanTM, represents the different components of a scientific project as top-level branches. These include a branch for raw data, along with separate branches for results, analysis, and publications related to the project. In the case of raw data, the semantics are not clear and cannot be defined by the platforms collecting the data. Scientists are able to sparsely attach meaning using horizontal links between different branches, using the OWL Open World Assumption to allow extra semantic annotations to be attached in future. For example, they may annotate images of a plant with a trait in only one case, without annotating other images. These horizontal links are not constrained by the basic scientific project structure, so they are able to entail novel meanings independent of the ontology, including untested hypotheses, process annotations and meaningful scientific results.

2.3 Semantic validation

PODD validates experiments using independently configurable rules based on OWL (Web Ontology Language) schemas. Although PODD currently only supports OWL as a rules language, it could be easily extended in other cases to use different systems such as N3, RDFS, SPARQL, or SPIN as rules languages [5].

OWL is used to determine whether projects are both internally consistent, with all objects having an explicit RDF type, and whether they are consistent with the ontologies that they import. The main points that PODD focuses on for validation is determining that mappings are consistent with the basic PODD predicates and classes, so that they can be used to render both static HTML pages and provide HTML based editing facilities. For example, any OWL object property that has been defined to link from image acquisition runs to images can be used to define the purpose of an image. This may include any combination of the defined properties, without requiring all of the properties to be used.

The parent object types and properties used by PODD to establish basic ground rules for rendering and validation are defined in a small base schema ontology. General scientific properties and phenotype specific properties are then defined in larger extension ontologies as illustrated in Figure 2.

An ontology is a common controlled vocabulary about a particular topic, in this case plant phenomics. The PODD Plant ontology therefore becomes an ontology for plant phenomics that extends the generally applicable PODD Science



Fig. 2. PODD ontology hierarchy

ontology with plant phenomics specific details [1]. This approach allows PODD to model multiple phenomics processes without enforcing a rigid data structure on objects, either now, in the past, or in the future. In addition, other community ontologies, such as the plant and gene ontologies [3, 6], can be used to annotate the data. This not only supports data discovery, but provides common reference points through which PODD managed data can be integrated with data in other research databases.

The PODD application handles the instantiation and validation of objects based on the versions of ontologies that are specified by the user as being relevant. This gives PODD a unique extensibility as the domain knowledge is entirely captured in the ontology, and the operational aspects of the system are derived purely from the content of the ontology. In particular, PODD is designed around the ability to support multiple versions of ontologies concurrently. Previously published artifacts can be loaded without modification into PODD without affecting current studies that may be using newer, incompatible, versions of the schema ontologies.

PODD is able to manage high volume, high resolution, heterogeneous datasets using a tree based structure, which in the PODD Science ontology is based on the concept of a “Project” (<http://podd.plantphenomics.org.au/podd/>). Data references are inserted into projects using subproperties of the PODD “contains” predicate. When results are available they are inserted independently and cross-referenced with the data that they were derived from using a subproperty of the PODD “refers to” property. Each of the subproperties have constraints on the types of objects that are valid as sources and destinations, making it easy to both validate existing uses, and suggest potential targets based on a source and a property.

The genotype of the plant alone does not determine the phenotype; it is the combination of the organisms genome with the environmental conditions that

determine the phenotype. PODD captures metadata on the genotype of the organism(s) under investigation. PODD also gathers comprehensive metadata on the environmental conditions and treatments that the biological samples are exposed to. This structure makes PODD suitable to queries by web agents thus linking information contained in other data repositories, e.g. genomics, transcriptomics, metabolomics to phenomics information.

3 The PlantScanTM digitisation platform

3.1 Heterogeneous data streams

PlantScanTM is a medium throughput high resolution phenotyping platform, which brings together a number of imaging sensors—light detection and ranging, far-infrared imaging, and multi-wavelength imaging—to non-invasively measure plant growth and function using in-silico approaches. Raw data is captured with its contextual information (e.g. system configuration, time of acquisition, batch number and project) and is stored in a purpose-built database as the data is being generated. The various data streams are collated and used to produce full 3D representation of each plant with overlaid spectral information. The metadata collected during image acquisition are necessary inputs for the computer vision techniques which are used to create the 3D representation of the plant. The 3D meshes are then automatically segmented in order to semantically identify the different parts of the plants (Paprocki et al, 2012). A longitudinal 3D matching pipeline for plant mesh parts is then used to evaluate temporal changes at the whole plant and/or organ level. After initial processing to recreate the 3D structure, the raw data for each plant is packaged into one location using BagIt, an hierarchical file [6]. All bags belonging to the same specific PODD project are grouped into collections on the CSIRO Data Access Portal. These bags have a permanent address which is recorded into PODD. Any additional processes on these digital objects (such as statistical information on features like plant colour, volume, growth etc. and on elements such as petiole, stem, leaf) are also captured by PODD (Figure 1). Metadata Each acquisition on PlantScanTM includes metadata (in addition to the raw data streams), such as plant genus and species, project and experiment metadata, a unique identifier for each image (Globally Unique Identifier), imaging angle, environmental temperature of the imaging chamber, location of optical and colour calibration datasets for each acquisition run, LiDAR (Light Detection and Ranging Sensors) calibration files. The metadata associated with each acquisition is automatically generated when setting up the configuration on the platform. This information is paramount to validate and process the raw image data, and for the post-processing phases. This information currently resides in the PlantScanTM database.

3.2 Data volume

Digitisation systems such as PlantScanTM generate huge amounts of data including raw image data, registration metadata, sensor configurations and plant

metadata. For example, PlantScanTM generates around 500GB of raw image data, representing in excess of 200,000 database records, per day. Sufficient storage space (usually at remote locations) and fast network transfer rates were thus necessary to facilitate data movement for processing using high performance computers (HPC). Because an RDF database structure is not suitable for handling large data sets of images, it was necessary to package the raw information into elementary units with permanent addresses which could be retrieved using PODD. The CSIRO Data Access Portal (DAP) [2] and Advanced Scientific Computing (ASC) storage and compute facilities [7] are key resources used by PlantScanTM to process and store bulk data.

4 CSIRO High Performance Compute Infrastructure

CSIRO ASC shared facilities [7] are used to process the raw PlantScanTM data to derive data products (meshes). Raw data and meshes are “bagged” and stored in the ASC archival system. ASC High Performance Compute (HPC) hosts (systems with high processor count and large memory) are taken advantage of to create and verify bags more rapidly than would be possible on conventional computer systems.

4.1 BagIt

BagIt is defined by an Internet Engineering Task Force (IETF) document as an “hierarchical file packaging format for storage and transfer of arbitrary digital content” [6]. A payload manifest details content and MD5 or SHA hashes for content integrity verification. Data file related metadata can be stored in predefined files as key-value pairs. For PlantScanTM, file-level metadata includes plant barcodes, batch numbers, and plant type.

The BagIt specification does not mandate a particular archiving strategy, the focus being upon the directory structure, special files, and integrity checking. Bag creation and payload manifest integrity checking, tools can take advantage of the multi-core capabilities of modern computing hardware. To suit current needs, BagIt-conforming tools [7] [8] [9] were assessed and where necessary, improvements were implemented and tested. For example, the implementation of [8] assumed that as many threads as possible should be created for multi-threaded operations such as bag creation and payload manifest verification. On a HPC host with hundreds of processors, this can lead to runaway thread creation, so this needed to be addressed. Neither [7] nor [8] permitted payload files in a bag to be listed. This has since been added to [8]. Some functionality (not discussed here) present in [8] but not in [7] was added to the latter.

5 CSIRO Data Access Portal

CSIRO’s Research Data Service (RDS) has developed the Data Access Portal (DAP), an open source web application that enables research data to be discovered, managed and shared. [5]

Researchers can describe a data collection, deposit data, choose a license, and add attribution details. Access to a collections description and/or data can be restricted to CSIRO or a set of individuals (within CSIRO or partner organisations) or it can be made public, becoming searchable by anyone via the Internet. In the case where a collection and its data are public, a Digital Object Identifier (DOI) is issued and can be used to formally cite the collection in a publication. Collections can either be one-time, self-serve deposits or automated, ongoing deposits and can be changed over time through versioning. Whole datasets or specific files within a collection can be downloaded. A machine-to-machine REST interface permits programmatic search and collection download.

DAP is the visible part of the system. Other than standard web servers, databases and application servers, additional required server-side infrastructure includes petabyte-level, multi-decade, backed-up disk storage, integrity checking, and automated migration of data collections to and from tape storage.

5.1 Bag preparation for a DAP collection

CSIROs HRPPC makes use of DAP to store collections of PlantScanTM raw images and processed mesh data as bags [6]. Raw data from PlantScanTM local storage (HRPPC-Store) and data processed on HPC hosts are transferred to ASC bulk storage where image and mesh files are organised in folders by batch, then barcode number, then subfolders for each image file type (e.g. RGB images, IR images, LiDAR) and meshes. Bag creation is carried out via an allocated ASC HPC job.

Currently, one bag is equivalent to a single batch scanned on the PlantScanTM local software system, which usually means the same kind of plant with different genotypes scanned under one experiment configuration profile. However the naming, content and amount of data in a bag are flexible and depend upon the intention of the bag. For example, multiple batches can be packed into one bag when the batches come from one plants scans at different time points. The multi-threaded BagIt tool is called on the ASC HPC host to create a bag that is then compressed into a single file containing a top-level data folder, a manifest file that contains checksums for every data file, and tag files that contain the metadata description of the bag just packed.

The metadata required for a DAP publication is created and the bag transferred to the DAP staging area via SFTP (SSH File Transfer Protocol). After publication of the DAP collection, the data from PlantScanTM for the given project becomes discoverable via DAP. In addition, experiment reports, published papers, and sensor configurations can either be made accessible via a DAP collection’s “related materials” links, other metadata fields, or within the collection’s data (e.g. bag).

6 Semantic integration

The PODD ontology enables plant phenomics researchers to link from mesh results to the raw data that they were generated from. It also allows researchers to

link from both mesh results and their recorded conclusions to shared phenomics ontologies which describe specific features of the plants. When used together, this enables scientists to trace the provenance of their results and conclusions based on well known concepts in phenomics ontologies.

The PODD ontology allows sparse annotation of data and results, so scientists are not required to annotate all of their data with links to shared ontologies by default. However, the system is extensible, so in future such a requirement could be added in a new version of the PODD ontology while maintaining backwards compatibility with previous experiments that may be published or immutable.

External ontologies must currently be mapped into the PODD system to define the expected status of the classes and properties with respect to the basic PODD concepts. The basic concepts are used to both verify that each PODD project has a single valid tree structure, and render HTML pages for both browsing and editing. PODD also supports an RDF-based HTTP REST interface that clients may use to fetch and update objects without using the HTML pages.

7 Semantic publication

PODD provides a secure mechanism for publishing both human and machine readable descriptions of scientific experiments. It utilises the well-known DOI mechanism for publishing raw data files using DAP, and uses HTTP URIs to publish experiments using the PODD web interface.

PODD implements support for Content Negotiation, with common RDF serialisations supported in addition to HTML pages that are encoded with RDFa (RDF in Attributes). Although most non-human agents will receive an RDF formatted document when they resolve a PODD URI, if a PODD HTML page is saved the annotations can still be retrieved using RDFa.

Scientific journals increasingly require the data and provenance for articles to be available in a machine readable format. The DOI registrar that DAP uses, DataCite [8], was setup to provide unique identifiers for data items that can be attached to publications, which in turn may have their own DOIs.

By providing machine readable descriptions of scientific experiments, including semantic references to shared ontologies where possible, PODD enables the output from PlantScanTM to be interpreted and extended by others. The use of PODD URIs in other RDF documents enables scientists to extend the initial work using the Linked Data paradigm [9].

8 Conclusion

This paper described how the Phenomics Ontology Driven Data repository integrates with the PlantScanTM platform and CSIRO Data Access Portal to manage the complex workflows at the High Resolution Plant Phenomics Centre. This

workflow keeps track of phenomics data, metadata and experimental processes and also provides a secure mechanism to share and publish scientific experiments in both human and machine readable formats.

References

1. Li, Y.F., Kennedy, G., Davies, F., Hunter, J.: Podd: An ontology-driven data repository for collaborative phenomics research. In Chowdhury, G., Koo, C., Hunter, J., eds.: The Role of Digital Libraries in a Time of Global Change. Volume 6102 of Lecture Notes in Computer Science. Springer Berlin Heidelberg (2010) 179–188
2. CSIRO IMT: Csiro data access portal. <http://data.csiro.au>
3. Ashburner, M., Ball, C.A., Blake, J.A., Botstein, D., Butler, H., Cherry, J.M., Davis, A.P., Dolinski, K., Dwight, S.S., Eppig, J.T., Harris, M.A., Hill, D.P., Issel-Tarver, L., Kasarskis, A., Lewis, S., Matese, J.C., Richardson, J.E., Ringwald, M., Rubin, G.M., Sherlock, G.: Gene ontology: tool for the unification of biology. the gene ontology consortium. *Nature Genet.* **25** (2000) 25–29
4. Belleau, F., Nolin, M.A., Tourigny, N., Rigault, P., Morissette, J.: Bio2rdf: Towards a mashup to build bioinformatics knowledge systems. *Journal of Biomedical Informatics* **41**(5) (2008) 706–716
5. Frber, C., Hepp, M.: Using sparql and spin for data quality management on the semantic web. In Abramowicz, W., Tolksdorf, R., eds.: Business Information Systems. Volume 47 of Lecture Notes in Business Information Processing. Springer Berlin Heidelberg (2010) 35–46
6. Avraham, S., Tung, C.W., Ilic, K., Jaiswal, P., Kellogg, E.A., McCouch, S., Pujar, A., Reiser, L., Rhee, S.Y., Sachs, M.M., Schaeffer, M., Stein, L., Stevens, P., Vincent, L., Zapata, F., Ware, D.: The plant ontology database: a community resource for plant structure and developmental stages controlled vocabulary and annotations. *Nucleic Acids Research* **36**(suppl 1) (2008) D449–D454
7. CSIRO IMT: Csiro advanced scientific computing. <https://wiki.csiro.au/display/ASC>
8. Brase, J.: Datacite - a global registration agency for research data. In: Cooperation and Promotion of Information Resources in Science and Technology, 2009. COINFO '09. Fourth International Conference on. (2009) 257–261
9. Berners-Lee, T. <http://www.w3.org/DesignIssues/LinkedData.html> (2006)

References [5] <https://wiki.csiro.au/confluence/display/dmsdoc/Home> [6] <http://datatracker.ietf.org/doc/dkunze-bagit/> [7] <https://pypi.python.org/pypi/bagit/> [8] <http://sourceforge.net/projects/loc-xferutils/> [9] <http://sourceforge.net/projects/loc-xferutils/files/loc-bagger/>